Seasonal Difference Exploration

Hao Zheng (hz2770)

5/6/2022

Explore the seasonal differences

```
load("./dt_long.RData")
load("./ID_in.RData")
load("./beta.res.postmean.RData")

dt_season <-
  dt_long %>%
  drop_na() %>%
  filter(ID %in% ID_in) %>%
  distinct(ID, .keep_all = TRUE) %>%
  select(ID, Season, Month, Nature) %>%
  mutate(Month = factor(Month, levels = month.name))
```

```
season_diff <-
  merge(dt_season, beta.res.postmean, by = c("ID")) %>%
  janitor::clean_names()
colnames(season_diff)[2] <- "year"

# Beta0
intercept.fit <- lm(intercept ~ month + year + nature, data = season_diff)
# Beta1
wind_prev.fit <- lm(wind_prev ~ month + year + nature, data = season_diff)
# Beta2
lat_change.fit <- lm(lat_change ~ month + year + nature, data = season_diff)
# Beta3
long_change.fit <- lm(long_change ~ month + year + nature, data = season_diff)
#Beta4
wind_change.fit <- lm(wind_change ~ month + year + nature, data = season_diff)

summary(intercept.fit)
```

```
##
## Call:
## lm(formula = intercept ~ month + year + nature, data = season_diff)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.43433 -0.03822  0.00234  0.04052  0.38709
##
```

```
## Coefficients:
##                  Estimate Std. Error t value Pr(>|t|)
## (Intercept)     4.4810021  0.3902677  11.482   <2e-16 ***
## monthApril      0.0232609  0.1113880   0.209   0.8346
## monthMay        0.0259813  0.0942017   0.276   0.7828
## monthJune       0.0275693  0.0922175   0.299   0.7651
## monthJuly       0.0125400  0.0918533   0.137   0.8914
## monthAugust    -0.0198034  0.0913669  -0.217   0.8285
## monthSeptember -0.0070528  0.0912856  -0.077   0.9384
## monthOctober    0.0093435  0.0913761   0.102   0.9186
## monthNovember   0.0145692  0.0924416   0.158   0.8748
## monthDecember   0.0057977  0.0976110   0.059   0.9527
## year           -0.0003419  0.0001895  -1.804   0.0717 .
## natureET        0.0008449  0.0298315   0.028   0.9774
## natureNR        0.0008122  0.0484835   0.017   0.9866
## natureSS        0.0141564  0.0205164   0.690   0.4904
## natureTS        0.0118370  0.0166932   0.709   0.4785
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.09022 on 682 degrees of freedom
## Multiple R-squared:  0.03095,    Adjusted R-squared:  0.01105
## F-statistic: 1.556 on 14 and 682 DF,  p-value: 0.0866
```

```
summary(wind_prev.fit)
```

```
##
## Call:
## lm(formula = wind_prev ~ month + year + nature, data = season_diff)
##
## Residuals:
##       Min        1Q    Median        3Q       Max
## -0.135790 -0.019433  0.002041  0.022272  0.061680
##
## Coefficients:
##                  Estimate Std. Error t value Pr(>|t|)
## (Intercept)     1.343e+00  1.215e-01  11.058  < 2e-16 ***
## monthApril      1.479e-02  3.467e-02   0.427 0.669679
## monthMay       -1.180e-04  2.932e-02  -0.004 0.996789
## monthJune       5.393e-03  2.870e-02   0.188 0.850987
## monthJuly       1.540e-02  2.859e-02   0.539 0.590174
## monthAugust     2.332e-02  2.843e-02   0.820 0.412418
## monthSeptember  2.610e-02  2.841e-02   0.919 0.358560
## monthOctober    2.108e-02  2.844e-02   0.741 0.458718
## monthNovember   2.461e-02  2.877e-02   0.856 0.392526
## monthDecember   8.824e-03  3.038e-02   0.290 0.771531
## year           -2.252e-04  5.899e-05  -3.817 0.000147 ***
## natureET        3.733e-03  9.284e-03   0.402 0.687709
## natureNR       -1.461e-02  1.509e-02  -0.969 0.333111
## natureSS       -3.330e-03  6.385e-03  -0.522 0.602172
## natureTS       -5.998e-03  5.195e-03  -1.155 0.248693
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 0.02808 on 682 degrees of freedom
## Multiple R-squared:  0.0659, Adjusted R-squared:  0.04672
## F-statistic: 3.437 on 14 and 682 DF,  p-value: 1.998e-05
```

```
summary(lat_change.fit)
```

```
##
## Call:
## lm(formula = lat_change ~ month + year + nature, data = season_diff)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.90321 -0.07062  0.00781  0.07691  0.95935
##
## Coefficients:
##                  Estimate Std. Error t value Pr(>|t|)
## (Intercept)     4.131e-02  6.667e-01   0.062    0.951
## monthApril      1.656e-02  1.903e-01   0.087    0.931
## monthMay        7.088e-02  1.609e-01   0.440    0.660
## monthJune      -7.088e-03  1.575e-01  -0.045    0.964
## monthJuly      -9.091e-03  1.569e-01  -0.058    0.954
## monthAugust    -5.225e-02  1.561e-01  -0.335    0.738
## monthSeptember -3.611e-02  1.559e-01  -0.232    0.817
## monthOctober   -2.862e-02  1.561e-01  -0.183    0.855
## monthNovember   2.400e-02  1.579e-01   0.152    0.879
## monthDecember  -5.431e-02  1.668e-01  -0.326    0.745
## year            3.655e-05  3.238e-04   0.113    0.910
## natureET       -7.020e-02  5.096e-02  -1.378    0.169
## natureNR        5.897e-03  8.283e-02   0.071    0.943
## natureSS       -1.352e-03  3.505e-02  -0.039    0.969
## natureTS       -1.545e-02  2.852e-02  -0.542    0.588
##
## Residual standard error: 0.1541 on 682 degrees of freedom
## Multiple R-squared:  0.02561,    Adjusted R-squared:  0.005609
## F-statistic:  1.28 on 14 and 682 DF,  p-value: 0.2137
```

```
summary(long_change.fit)
```

```
##
## Call:
## lm(formula = long_change ~ month + year + nature, data = season_diff)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.30817 -0.03599  0.00530  0.04273  0.50782
##
## Coefficients:
##                  Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -0.8336700  0.3531651  -2.361   0.0185 *
## monthApril      0.0416468  0.1007984   0.413   0.6796
## monthMay        0.0632772  0.0852459   0.742   0.4582
## monthJune       0.0556884  0.0834504   0.667   0.5048
## monthJuly       0.0361214  0.0831208   0.435   0.6640
```

```
## monthAugust     0.0123691  0.0826807   0.150   0.8811
## monthSeptember  0.0212965  0.0826071   0.258   0.7966
## monthOctober    0.0341549  0.0826890   0.413   0.6797
## monthNovember   0.0263450  0.0836532   0.315   0.7529
## monthDecember   0.0422468  0.0883312   0.478   0.6326
## year            0.0002184  0.0001715   1.273   0.2033
## natureET       -0.0263888  0.0269955  -0.978   0.3287
## natureNR        0.0030556  0.0438742   0.070   0.9445
## natureSS        0.0126339  0.0185659   0.680   0.4964
## natureTS       -0.0231521  0.0151062  -1.533   0.1258
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.08164 on 682 degrees of freedom
## Multiple R-squared:  0.05042,    Adjusted R-squared:  0.03093
## F-statistic: 2.586 on 14 and 682 DF,  p-value: 0.001201
```

```
summary(wind_change.fit)
```

```
##
## Call:
## lm(formula = wind_change ~ month + year + nature, data = season_diff)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.40181 -0.04476 -0.00309  0.04544  0.35691
##
## Coefficients:
##                  Estimate Std. Error t value Pr(>|t|)
## (Intercept)     2.890e-01  3.809e-01   0.759    0.448
## monthApril      3.618e-02  1.087e-01   0.333    0.739
## monthMay       -1.629e-02  9.195e-02  -0.177    0.859
## monthJune       2.377e-02  9.001e-02   0.264    0.792
## monthJuly       1.308e-02  8.965e-02   0.146    0.884
## monthAugust     3.124e-02  8.918e-02   0.350    0.726
## monthSeptember  4.448e-02  8.910e-02   0.499    0.618
## monthOctober    3.505e-02  8.919e-02   0.393    0.694
## monthNovember   2.091e-02  9.023e-02   0.232    0.817
## monthDecember   1.142e-02  9.527e-02   0.120    0.905
## year            9.048e-05  1.850e-04   0.489    0.625
## natureET       -2.092e-02  2.912e-02  -0.719    0.473
## natureNR       -2.173e-02  4.732e-02  -0.459    0.646
## natureSS       -2.385e-02  2.003e-02  -1.191    0.234
## natureTS       -1.750e-02  1.629e-02  -1.074    0.283
##
## Residual standard error: 0.08806 on 682 degrees of freedom
## Multiple R-squared:  0.02104,    Adjusted R-squared:  0.000943
## F-statistic: 1.047 on 14 and 682 DF,  p-value: 0.404
```

```
sum0 <- summary(intercept.fit)$coefficients[,c(1,4)]
sum1 <- summary(wind_prev.fit)$coefficients[,c(1,4)]
sum2 <- summary(lat_change.fit)$coefficients[,c(1,4)]
sum3 <- summary(long_change.fit)$coefficients[,c(1,4)]
```

	Beta 0		Beta 1		Beta 2		Beta 3	
	Estimate	Pr(>|t|)	Estimate	Pr(>|t|)	Estimate	Pr(>|t|)	Estimate	Pr(>|t|)
(Intercept)	4.4810021	0.0000000	1.3431063	0.0000000	0.0413063	0.9506172	-0.8336700	0.0185275
monthApril	0.0232609	0.8346449	0.0147943	0.6696787	0.0165579	0.9306863	0.0416468	0.6796126
monthMay	0.0259813	0.7827813	-0.0001180	0.9967888	0.0708822	0.6597505	0.0632772	0.4581672
monthJune	0.0275693	0.7650618	0.0053935	0.8509869	-0.0070875	0.9641298	0.0556884	0.5047909
monthJuly	0.0125400	0.8914489	0.0154032	0.5901741	-0.0090910	0.9538180	0.0361214	0.6640154
monthAugust	-0.0198034	0.8284715	0.0233206	0.4124181	-0.0522548	0.7378961	0.0123691	0.8811234
monthSeptember	-0.0070528	0.9384385	0.0261005	0.3585599	-0.0361073	0.8169707	0.0212965	0.7966351
monthOctober	0.0093435	0.9185853	0.0210829	0.4587183	-0.0286163	0.8546050	0.0341549	0.6796975
monthNovember	0.0145692	0.8748155	0.0246144	0.3925264	0.0239972	0.8792681	0.0263450	0.7529105
monthDecember	0.0057977	0.9526542	0.0088244	0.7715305	-0.0543131	0.7447475	0.0422468	0.6326060
year	-0.0003419	0.0717253	-0.0002252	0.0001471	0.0000365	0.9101708	0.0002184	0.2032812
natureET	0.0008449	0.9774141	0.0037334	0.6877086	-0.0702038	0.1687975	-0.0263888	0.3286540
natureNR	0.0008122	0.9866387	-0.0146142	0.3331114	0.0058967	0.9432660	0.0030556	0.9444979
natureSS	0.0141564	0.4904257	-0.0033299	0.6021721	-0.0013517	0.9692484	0.0126339	0.4964264
natureTS	0.0118370	0.4785102	-0.0059979	0.2486925	-0.0154533	0.5880814	-0.0231521	0.1258337

```
sum4 <- summary(wind_change.fit)$coefficients[,c(1,4)]

kable(cbind(sum0, sum1, sum2, sum3, sum4)) %>%
  kable_styling(bootstrap_options = c("striped", "hover", "condensed")) %>%
  add_header_above(c(" " = 1, "Beta 0" = 2, "Beta 1" = 2, "Beta 2" = 2, "Beta 3" = 2, "Beta 4" = 2))
```

Now based on the estimated Bayesion model from previous questions, we need to explore the seasonal difference. We can fit 5 models using 5 estimated beta values against the three predictors: $X_{i,1}$: the month of the year the ith hurricane started, $X_{i,2}$:the year of the ith hurricane and $X_{i,3}$: the nature of the ith hurricane. The beta values obtained from previous Gibbs Sampler MCMC method contains the mean value of $\beta_{0,i}$, $\beta_{1,i}$, $\beta_{2,i}$, $\beta_{3,i}$ and $\beta_{4,i}$ for each of the 697 unique hurricanes, which is of the size 697 * 5.

According to the summary, the R squared value for all the five fitted linear models are quite small, which may indicate bad fit. In addition, most coefficients for the model are not significant with a p-value larger than 0.05. However, for those significant coefficients, we could infer a potential relationship between the certain predictors and the beta coefficients respectively. We should consult the previous Bayesion model:

$$Y_i(t+6) = \beta_{0,i} + \beta_{1,i}Y_i(t) + \beta_{2,i}\Delta_{i,1}(t) + \beta_{3,i}\Delta_{i,2}(t) + \beta_{4,i}\Delta_{i,3}(t) + \epsilon_i(t)$$

to interpret the change of the influence on $Y_{i,t+6}$ as the value of the predictor changes.

For the fitted coefficients of β_0 to β_4, the intercept cannot show information about seasonal difference since they indicate when holding all the predictors zero, the value for the corresponding β. We can only observe that the year is quite significant in the model for β_0, β_1 with both negative estimates close to zero. Therefore, as the year increase, the coefficient of the intercept and $Y_{i,t}$ may decrease a little, which means for the Bayesian model, the wind speed when holding all the variables zero and the effect the previous wind speed has will decrease over years. Apart from seasonal difference, some other predictors are quite significant, such as natureET for β_2, natureTS for β_3.

```
# Try to fit the beta model only with the four seasons
season_diff <- as.data.frame(season_diff) %>%
  mutate(month = recode(month, April = "Spring"),
       month = recode(month, May = "Spring"),
       month = recode(month, June = "Summer"),
```

	Beta 0		Beta 1		Beta 2		Beta 3		
	Estimate	Pr(>|t|)	Estimate	Pr(>|t|)	Estimate	Pr(>|t|)	Estimate	Pr(>|t|)	E
(Intercept)	3.8365500	0.0000000	0.8942250	0.0000000	0.1606506	0.0000610	-0.3500900	0.0000000	0.
seasonSummer	-0.0305003	0.2048954	0.0152377	0.0440074	-0.0979486	0.0167511	-0.0466127	0.0338037	0.
seasonAutumn	-0.0235346	0.3248438	0.0209616	0.0053662	-0.0909590	0.0253577	-0.0434764	0.0463302	0.
seasonWinter	-0.0186542	0.6535827	0.0034158	0.7936540	-0.0984181	0.1637856	-0.0094850	0.8023902	0.

```
       month = recode(month, July = "Summer"),
       month = recode(month, August = "Summer"),
       month = recode(month, September = "Autumn"),
       month = recode(month, October = "Autumn"),
       month = recode(month, November = "Autumn"),
       month = recode(month, December = "Winter"),
       month = recode(month, January = "Winter"),
       month = factor(month, levels = c("Spring", "Summer", "Autumn", "Winter")))
colnames(season_diff)[3] <- "season"

# Beta0
intercept.fit.2 <- lm(intercept ~ season, data = season_diff)
# Beta1
wind_prev.fit.2 <- lm(wind_prev ~ season, data = season_diff)
# Beta2
lat_change.fit.2 <- lm(lat_change ~ season, data = season_diff)
# Beta3
long_change.fit.2 <- lm(long_change ~ season, data = season_diff)
# Beta4
wind_change.fit.2 <- lm(wind_change ~ season, data = season_diff)

sum0_2 <- summary(intercept.fit.2)$coefficients[,c(1,4)]
sum1_2 <- summary(wind_prev.fit.2)$coefficients[,c(1,4)]
sum2_2 <- summary(lat_change.fit.2)$coefficients[,c(1,4)]
sum3_2 <- summary(long_change.fit.2)$coefficients[,c(1,4)]
sum4_2 <- summary(wind_change.fit.2)$coefficients[,c(1,4)]

kable(cbind(sum0_2, sum1_2, sum2_2, sum3_2, sum4_2)) %>%
  kable_styling(bootstrap_options = c("striped", "hover", "condensed")) %>%
  add_header_above(c(" " = 1, "Beta 0" = 2, "Beta 1" = 2, "Beta 2" = 2, "Beta 3" = 2, "Beta 4" = 2))
```

We also try to represent the months as four seasons and fit a model for β against them. Each model has three dummy variables corresponding to the three seasons except Spring. The latter three rows of estimate shows how the value of β differentiate between Spring and the other three seasons respectively. If with a rather small p-value, we can conclude the existence of seasonal difference. Therefore, by constructing model in this way, we find that β_1 and β_4 will increase a little as season changes from Spring to Summer, then to Autumn, which means a season difference of the effect $Y_{i,t}$ and $\Delta_{i,3}(t)$ has on the wind speed. For β_2, β_3, Summer and Autumn may lead to a slightly smaller effect of $\Delta_{i,1}(t)$, $\Delta_{i,2}(t)$ have on the wind speed compared to Spring.

```
# Try to fit the beta model only with the year
# Beta0
intercept.fit.new <- lm(intercept ~ year, data = season_diff)
# Beta1
wind_prev.fit.new <- lm(wind_prev ~ year, data = season_diff)
```

```r
# Beta2
lat_change.fit.new <- lm(lat_change ~ year, data = season_diff)
# Beta3
long_change.fit.new <- lm(long_change ~ year, data = season_diff)
#Beta4
wind_change.fit.new <- lm(wind_change ~ year, data = season_diff)

summary(intercept.fit.new)
```

```
##
## Call:
## lm(formula = intercept ~ year, data = season_diff)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.43803 -0.03487  0.00150  0.04113  0.37030
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4.5142875  0.3580505  12.608   <2e-16 ***
## year        -0.0003543  0.0001803  -1.965   0.0498 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.09053 on 695 degrees of freedom
## Multiple R-squared:  0.005526,   Adjusted R-squared:  0.004095
## F-statistic: 3.862 on 1 and 695 DF,  p-value: 0.04979
```

```r
summary(wind_prev.fit.new)
```

```
##
## Call:
## lm(formula = wind_prev ~ year, data = season_diff)
##
## Residuals:
##       Min        1Q    Median        3Q       Max
## -0.131834 -0.019873  0.000843  0.023479  0.060412
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.3448481  0.1126207  11.941  < 2e-16 ***
## year        -0.0002178  0.0000567  -3.842 0.000133 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.02848 on 695 degrees of freedom
## Multiple R-squared:  0.0208, Adjusted R-squared:  0.01939
## F-statistic: 14.76 on 1 and 695 DF,  p-value: 0.0001332
```

```r
summary(lat_change.fit.new)
```

```
##
```

```
## Call:
## lm(formula = lat_change ~ year, data = season_diff)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.90041 -0.06349  0.00594  0.07630  1.07461
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.056e-01  6.117e-01  -0.173    0.863
## year         8.776e-05  3.079e-04   0.285    0.776
##
## Residual standard error: 0.1547 on 695 degrees of freedom
## Multiple R-squared:  0.0001168,  Adjusted R-squared:  -0.001322
## F-statistic: 0.08122 on 1 and 695 DF,  p-value: 0.7757
```

```
summary(long_change.fit.new)
```

```
##
## Call:
## lm(formula = long_change ~ year, data = season_diff)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.31082 -0.02940  0.00746  0.04013  0.57851
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.0267628  0.3273439  -3.137  0.00178 **
## year         0.0003188  0.0001648   1.934  0.05347 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.08277 on 695 degrees of freedom
## Multiple R-squared:  0.005355,   Adjusted R-squared:  0.003924
## F-statistic: 3.742 on 1 and 695 DF,  p-value: 0.05347
```

```
summary(wind_change.fit.new)
```

```
##
## Call:
## lm(formula = wind_change ~ year, data = season_diff)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.40370 -0.04331 -0.00361  0.04853  0.36826
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 3.051e-01  3.486e-01   0.875    0.382
## year        9.024e-05  1.755e-04   0.514    0.607
##
## Residual standard error: 0.08814 on 695 degrees of freedom
```

	Beta 0		Beta 1		Beta 2		Beta 3		
	Estimate	Pr(>|t|)	Estimate	Pr(>|t|)	Estimate	Pr(>|t|)	Estimate	Pr(>|t|)	Estim
(Intercept)	4.5142875	0.0000000	1.3448481	0.0000000	-0.1056332	0.8629385	-1.0267628	0.001781	0.3051
year	-0.0003543	0.0497902	-0.0002178	0.0001332	0.0000878	0.7757368	0.0003188	0.053474	0.0000

```
## Multiple R-squared:  0.0003802,  Adjusted R-squared:  -0.001058
## F-statistic: 0.2644 on 1 and 695 DF,  p-value: 0.6073
```

```
sum0.new <- summary(intercept.fit.new)$coefficients[,c(1,4)]
sum1.new <- summary(wind_prev.fit.new)$coefficients[,c(1,4)]
sum2.new <- summary(lat_change.fit.new)$coefficients[,c(1,4)]
sum3.new <- summary(long_change.fit.new)$coefficients[,c(1,4)]
sum4.new <- summary(wind_change.fit.new)$coefficients[,c(1,4)]

kable(cbind(sum0.new, sum1.new, sum2.new, sum3.new, sum4.new)) %>%
  kable_styling(bootstrap_options = c("striped", "hover", "condensed")) %>%
  add_header_above(c(" " = 1, "Beta 0" = 2, "Beta 1" = 2, "Beta 2" = 2, "Beta 3" = 2, "Beta 4" = 2))
```

Now fit linear models for β against the season variables (corresponding to the year) to seek for potential evidence of the statement :"the wind speed has been increasing over years". In order to analyze this question, need to inspect on model which corresponds to the wind speed and the year. For β_2 model, the estimate of year is significant, although it's really close to zero. Therefore, we can infer that as the year increases, the impact past wind speed has on the current wind speed may decrease a little, which cannot provide support for the statement. However, it's quite match with the results shown in the figures in the initial EDA session, which indicates the mean wind speed tends to decrease over years.

In conclusion, for different months, there is no significant differences observed. Over years, the effect the wind speed 6 months ago has on the current wind speed may decrease a little. And there is no evidence to support the statement in task 5.